

**CPU-I/O Cycle:** Process execution consists of a cycle of:

1. CPU execution (**CPU burst time**)
2. I/O wait (**I/O burst time**)

**CPU Scheduler:** Selects a process from the ready queue

- **Non-pre-emptive:** Process runs until it voluntarily relinquishes CPU
- **Preemptive:** Process runs for a maximum of some fixed time

**Average Waiting time calculation:** Calculate wait time for each process and then average it

**Scheduling Algorithms:**

- **FCFS:** First come first serve
  - o **Problem:** ???
- **LIFO:** Last in first out
  - o **Problem:** May lead to starvation – early processes may never get CPU
- **SJF:** Shortest job first
  - o Estimated CPU burst time is associated with each process. Scheduler uses these lengths to schedule the process with the shortest time
  - o Optimal – gives minimum average waiting time for a given set of processes
- **Priority Scheduling:** A priority number is associated with each process
  - o **Problem:** Starvation
  - o **Solution:** Aging
    - As time progresses, increase the priority of the process
- **RR:** Round Robin
  - o Each process gets a small unit of CPU time (**time quantum**). After this time has elapsed, the process is pre-empted and added to the end of the ready queue
  - o If there are  $n$  processes in the ready queue, and the time quantum is  $q$ , then each process gets at most  $q$  time units at once. No process waits for more than  $(n-1)q$  time units
  - o If  $q$  is too large: FIFO-like behaviour
  - o If  $q$  is too small: context switching

**Multilevel Queue Scheduling:** Uses multiple queues. Essentially the ready queue is really multiple separate queues. The processes can be classified into different groups (i.e. foreground vs background)

- Each queue has its own scheduling algorithm
- Scheduling must be done between the queues:
  - o **Fixed priority Scheduling:** Serve all from foreground then from background
  - o **Time Slice:** Each queue gets a certain amount of CPU time which it can schedule amongst its processes
- A process can move between queues. Separate processes according to the characteristics of the CPU bursts (i.e. if a process has too much CPU time, it will be moved to a lower

priority queue, or a process that waits too long in a low priority queue may be moved to a higher priority one)

## Multiple-Processor Scheduling

- **Asymmetric Multiprocessing:** There is one processor that makes the decisions for scheduling, I/O processing, and system activities
  - o Other processor(s) execute only user code
  - o Simple approach to master-slave model / centralized command model
- **Symmetric Multiprocessing:** Each processor is self-scheduling, and all processors share a **common ready queue** OR each processor may have its own **private queue** of ready processes

**Race Condition:** Behaviour of a system where the output is dependent on the sequence or timing of other uncontrollable events. It becomes a bug when events do not happen in the order the programmer intended

- Example: We have an application that withdraws money from the bank. Two requests for withdrawal from the same account comes to a bank from two different ATM machines
  - o Assume a balance of \$1000
  - o A thread for each request is created

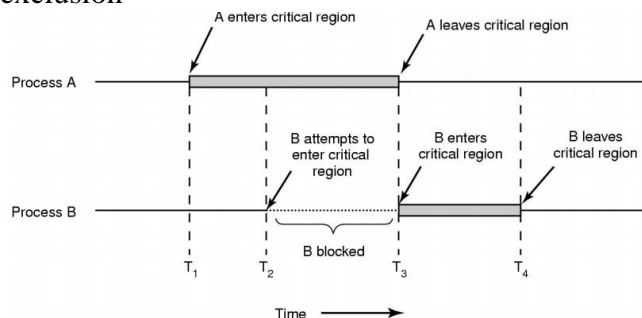
Process / Thread 1	Process / Thread 2
1. Read balance: \$1000	
2. Withdraw authorized for \$600 (now actual balance is \$400)	
CPU switches to process 2 →	3. Read Balance \$1000
	4. Withdraw authorized for \$600 (this is unreal!!)
5. Update balance \$1000-\$600 = \$400	← CPU switches to process 1
CPU switches to process 2 →	6. Update balance \$400-\$600 = \$-200

o

**Critical Section:** Any piece of code that accesses shared data (i.e. printer, bank account)

**Mutual Exclusion:** Ensures that only one thread/process accesses the critical section at a time

- Today the assumption is that the OS will provide some sort of support for mutual exclusion



## Mutual Exclusion Solutions:

- **Peterson's Solution:** Restricted to 2 processors
  - 2 variables are shared, **int turn** and **int flag[2]**
  - The variable turn indicates whose turn it is to enter the critical section
  - The flag array is used to indicate if a process is ready to enter the critical section (flag[i] = true implies that process P<sub>i</sub> is ready)

Process 1	Process 2
flag[0] = false; flag[1] = false;	flag[0] = false; flag[1] = false;
{	{
flag[0] = true;	flag[1] = true;
turn = 2;	turn = 1;
while (flag[1] && turn == 2);	while (flag[0] && turn == 1);
critical section	critical section
flag[0] = false;	flag[1] = false;
remainder section	remainder section
}	}

- 
- **Disabling Interrupts:** Guarantees that there will be no process switch
  - Process disables all interrupts before entering its critical section
  - Process enables all interrupts just before leaving the critical region
  - CPU is switched from one process to another only via clock/interrupts

```
bank_example (account, amount_to_withdraw)
{
    disable(interrupts);
    1. balance = get_balance(account);

    2. if (balance == amount_to_withdraw)
        withdraw_authorized();
        else
            withdraw_request_denied()

    3. balance = balance - amount;

    enable(interrupts);
}
```

- 
- **Disadvantages:**
  - Gives the power to control interrupts to user
  - Does not work in the case of multiple CPUs. Only the CPU that executes the disable instruction is effected

- **Test and Lock Instruction (TSL):**

- Many computers have a TSL REGISTER, LOCK instruction
  - Reads LOCK into register REGISTER
  - Stores a nonzero value at the memory location LOCK
  - The CPU executing the TSL instruction locks the memory bus to prohibit other CPUs from accessing the memory until TSL instruction is done
- **Disadvantages:** Requires **busy waiting**, which wastes CPU cycles that some other process might be able to use

```

enter_region
TSL REGISTER, LOCK
CMP REGISTER, #0
JNE enter_region
Critical Section

```

- If lock value != 0
- JNE causes control to go to the start of enter\_region
- If lock value is zero then enter CS

○

- **Semaphores:** A process synchronization technique supported by the OS. Used to ensure mutual exclusion and send signals from one process to another

- Allows multiple processes to cooperate by using signal. Semaphore is an integer variable with the following three operations:
  - **Initialize:** Semaphore (S) is initialized to a positive value
  - **Decrement:** (down operation) decrements the semaphore value
  - **Increment:** (up operation) increments the semaphore value
- A process can enter a critical section only if S is positive
- Two types of semaphores:
  - **Binary:** Allows only ONE process to be in critical section at a time
    - Initialized to 1
    - Also referred to as **mutex**
  - **Counting:** Allows multiple processes to be in critical section at a time
    - Initialized to N where N is the maximum processes that can be in critical section simultaneously

- Use **down** before entering a critical section
- Use **up** after finishing with a critical section
- Example: Assume S is initialized to 1.

```

□   S = 1;
    while (true)
    {
        down(S);
        critical section
        up(S);
        remainder section;
    }

```

○

**Deadlock:** Two or more processes are waiting indefinitely for an event that can only be caused by one of the waiting processes

- **Conditions:**
  - Resource allocation cycle
  - Mutual exclusion
  - Hold and wait
  - Non pre-emptive resource
- **Avoidance:**
  - Avoid cycles in the resource allocation graph
  - Avoid mutual exclusion
  - Avoid hold & wait
  - Block a process that is requesting a large amount of resources
- **Recovery:**
  - Abort all deadlock processes
  - Back up all deadlock processes to the previous safe state and then restart
  - Selectively abort processes until deadlock is broken

**Memory Management:** Memory management requires

- Allocate memory to processes when needed
- Keep track of what parts of memory are in use
- Deallocate memory when processes are done

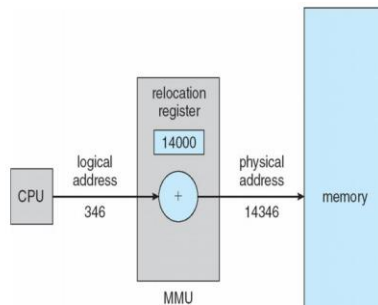
**Memory Hierarchy**

- Registers
- On-chip Cache (fast memory between the CPU and main memory so that CPUs don't have to wait for main memory)
- Main Memory
- Disk

**Address Binding:** Memory references in the code (**virtual** or **logical**) must be translated to actual physical memory addresses

**Memory Management Unit (MMU):** A hardware device that does run-time mapping from virtual to physical addresses

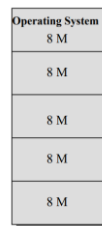
- The **base** register holds the physical memory address
- The **limit** register specifies the range
- These registers can be loaded only by the OS
- Ensures the user program doesn't access anything beyond its range
- **Relocation register:** Value is added to every address generated by a user process



## Memory Allocation Techniques

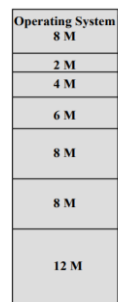
- **Contiguous Memory Allocation:** Each process is contained in a single section of memory that is contiguous
  - o **Fixed Partitioning:** Any program. No matter how small, occupies an entire partition

- Leads to **internal fragmentation**
- No modern OS uses fixed partitioning



- o **Dynamic Partitioning:** Partitions are a variable length and number. Processes are allocated to the closest possible match

- Leads to **external fragmentation**
- **Compaction** is required to obtain a large contiguous block
  - Shift processes so they are contiguous and all free memory is in one block



## Dynamic Partitioning Placement Algorithm

- Each process can be assigned to the smallest partition within which it will fit
- Processes are assigned in such a way as to minimize wasted memory within a partition
- Operating system must decide which free block to allocate to a process
- **Best-fit Algorithm:** Choose the block that is closest in size to the request
  - o This has worst overall performance
  - o The smallest block is found for a process
- **First-fit Algorithm:** Starts scanning memory from the beginning and chooses the first available block that is large enough

**Paging:** A memory management technique that avoids external fragmentation and compaction

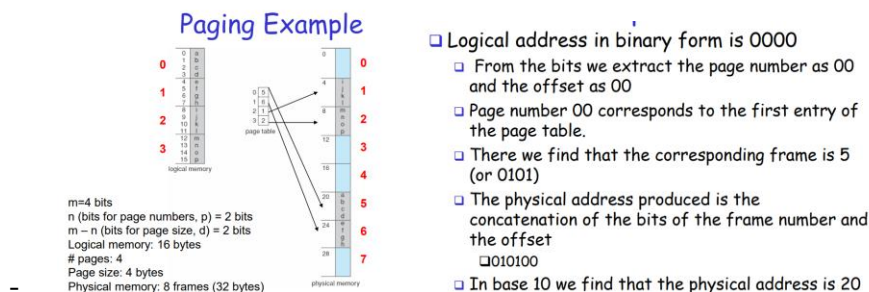
- Partition physical memory into small equal-size chunks and divide a process's logical memory into the same size chunks
- The chunks of a process are called **pages** and chunks of memory are called **frames**
- When a process is to be executed, its pages are loaded (from a file system) into any available memory frames
- A OS maintains a **page table** for each process that contains the frame location for each page of a process

**Address Translation Scheme:** Logical address generated by CPU is divided into:

- **Page number (p):** Used as an index to a page table which contains the frame number
- **Page offset (d):** Determines the location in the page which is needed to determine the location in the frame
- If the number of logical address bits is **m** and the number of bits to represent page numbers is **n**:
  - o Logical address space =  $2^m$
  - o Page numbers =  $2^n$
  - o Page size is  $2^{m-n}$

## Paging Hardware:

- The CPU issues a logical address
- The hardware extracts the page number **p** and the page offset **d**
- The page number **p** is used to index the page table (the entry of the page table consists of the frame number **f**)
- The actual address is the concatenation of the bits that make up **f** and **d**



**Page Table Implementation:** The simplest approach is to have the page table implemented as a set of dedicated registers. Not feasible to keep page tables in registers. Because page tables can be very large

- Each process has a page table
- Page table is kept in main memory
- **Page-table base register (PTBR)** points to the page table
- During a context switch, changing page tables requires changing PTBR
- Solution: Use a special fast-lookup hardware cache called associative memory or **translation look-aside buffers (TLBs)**

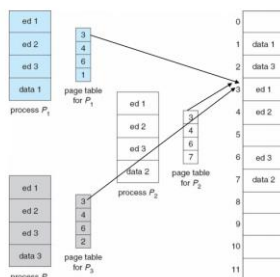
**Translation Look-Aside Buffers (TLB):** The TLB contains only a few of the page-table entries. When a logical to physical address is requested by the CPU its page number is presented to the TLB

- Address translation (p,d)
  - o If **p** is in associative memory, get frame number out
  - o Otherwise get frame number from page table in memory
- **TLB hit:** If found the frame number is immediately available
- **TLB miss:** If a page number is not in the TLB
  - o The page table is consulted
  - o The page number and frame number is added to the TLB
  - o If the TLB is full then one of the entries is replaced (discussed later)
- **Hit Ratio:** Percentage of time that a particular page is found
- **Miss Ratio:** Percentage of time that a particular page is not found (100 – hit ratio)
- **Effective Access Time:**  $EAT = \text{hit ratio} * \text{TLB hit time to get data} + \text{miss ratio} * \text{TLB miss time to get data}$ 
  - o i.e. let Hit ratio = 80%
  - o Therefore, Miss ratio = 20%
  - o TLB hit time to get data = 120
  - o TLB miss time to get data = 220
  - o Effective access time =  $0.80 * 120 + 0.20 * 220$

**Page Protection:** Memory protection implemented by associated protection bit with each frame

- One protection bit can define a page to be read-write or read-only

**Shared Pages:**

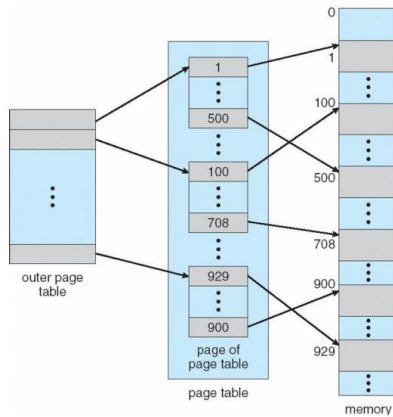




**Page Table Structure:** Typically systems have large logical address spaces

- Page table could be excessively large
- i.e. 32-bit logical address space,  $\therefore$  the number of possible address in the logical address space is  $2^{32}$
- Page table may consist of up to 1 million entries (very large)

**Two-Level Page Table:** Solution to large page tables. A page table is also paged. Need to be able to index the outer page table



**Virtual Memory:** This is basically logical address space. Processes use a virtual (logical) address space

- Every process has its own address space
- Only part of the virtual address space is mapped to physical memory at any time

**Demand Paging:** Bring a page into memory only when it is needed.

- Why? Less I/O needed and less memory needed
- We need to distinguish between pages that are in memory and the pages that are on disk
- A valid-invalid bit is part of each page entry
  - o When bit is 'valid', associated page is in memory
  - o When bit is 'invalid', the page is on the disk

**Page Fault:** Happens if a process tries to access a page that was not brought to memory

**Page Replacement:** Determine what page entry to replace when page table is full

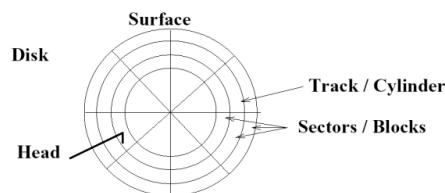
- Designing an appropriate algorithm is important since disk I/O is expensive
- **Optimal Page Replacement:** Replace page needed at the farthest point in future (i.e. replace the page that will not be used for the longest period of time)
  - o Lowest page fault rate
  - o Easy to describe but impossible to implement (OS has no way of knowing when each of the pages will be referenced next)
- **FIFO Page Replacement:** Maintain a linked list of all pages. Each page is associated with the time when that page was brought into memory. Page chosen to be replaced is the oldest page
  - o Implementation is a FIFO queue
    - Head points to the oldest page
    - Tail points to the newest page
- **Least Recently Used (LRU) Page Replacement:** Associates with each page the time of that page's last use. When a page must be replaced, LRU chooses the page that has not been used for the longest period of time
  - o Most OS's use LRU
- Other algorithms:
  - o Least Frequently Used (LFU)
  - o Most Frequently Used (MFU)

**Long-term Information Storage:** Three essential requirements:

- Must store large amounts of data
- Information stored must survive the termination of the process using it
- Multiple processes must be able to access the information concurrently

**Mass Storage Structures:**

- **Magnetic Disks:** Provide the bulk of secondary storage for modern computer systems



- **Tracks:** concentric rings on **platter** (see above)
    - o bits laid out serially on tracks
  - Tracks split into **sectors**
  - o **Layout:** □ Sectors may be grouped into blocks
- **Magnetic Tape:** Was used as an early secondary storage medium
  - o Slow compared to magnetic disks & memory
  - o Can hold large amounts of data
  - o Read data sequentially
  - o Mainly used for backup

**Disk Pack: Multiple Disks:** Think of disks as a stack of platters. Use both sides of platters (rather than 1 side)

- Two **read-write heads** at each end of arm
- **Cylinders:** Same track on each surface

**Disk Operations Cost:** Access time composed of:

- **Seek time:** Time to position head over correct track
- **Rotation time:** Time for correct sector to rotate under disk head
- **Transfer time:** Time to transfer data
- Usually seek time takes the longest

**I/O Bus:** A disk drive is attached to a computer by a set of wires called an I/O bus. Types of I/O bus include:

- Enhanced Integrated Drive Electronics (EDIE)
- Advanced Technology Attachment (ATA)
- Serial ATA (SATA)
- Universal Serial Bus (USB)
- Small Computer Systems Interface (SCI)

**I/O Controllers:** The data transfers on a bus are carried out by special electronic processors called controllers

- **Host controller:** Controller at the computer end of the bus
- **Disk controller:** Controller that is built into each disk drive

**Network Attached Storage (NAS):** Storage made available over a network rather than over a local connection

- Often a set of disks (storage array) are placed together
- NFS and CIFS are common protocols

**Disk Scheduling:** The disk accepts requests, what sort of disk arm scheduling algorithm is needed?

- **First Come First Serve (FCFS):**
  - o **Pros:** Fairness among requests
  - o **Cons:** Arrival may be on random spots on the disk (long seeks)
- **Shortest Seek Time First (SSTF):** Pick the one closest on disk
  - o **Pros:** Minimize seek time
  - o **Cons:** Starvation
  - o Most often used
- Others:
  - o SCAN / C-SCAN

**RAID:** Redundant Array of Independent Disks. Uses multiple disks attached to a computer system

- Provides reliability and redundancy
- Improve performance via **parallelism**
- Improve reliability via **information redundancy**
- Different scheme to provide redundancy, classified as **RAID Levels:**
  - o **RAID-0:** Break a file into blocks of data, stripe the blocks across disks in the system
    - Provides no redundancy or error detection
    - Loss of 1 disk means all data among both disks is lost
    - Allow to use 2 disks as 1. Used in systems where fast reads are required but minimal data integrity is needed
  - o **RAID-1:** A complete file is stored on a single disk, the second disk contains an exact copy of the file
    - Provides complete redundancy of data
    - Write performance suffers (must write data twice)
    - Most expensive implementation (need twice as much storage space)

**File Attributes:** Name, type, location, size, protection, creation time, etc

**File Naming:** Files with a certain standard structure imposed can be identified using an **extension** to their name

**File Operations:** Create, write, read, delete, etc

- For write/read operations, the OS needs to keep a **file position pointer** for each process

**Directory Operations:** File search, creating, deletion, renaming, directory listing

**File Systems:**

- Unix uses **Unix File System (UFS)**
- Windows **NT File System (NTFS)**, FAT, FAT32

**File Control Block (FCB):** Information about a file may be maintained in an FCB

- One FCB per file
- A FCB is associated with a unique id
- Consists of details about a file (permissions, dates, owners/groups, size, etc)

**File System Structures:**

- **System-wide Open File Table:** Contains a copy of the FCB of each open file
- **Per-process open-file table:** Contains a pointer to the appropriate entry in the system-wide open-file table

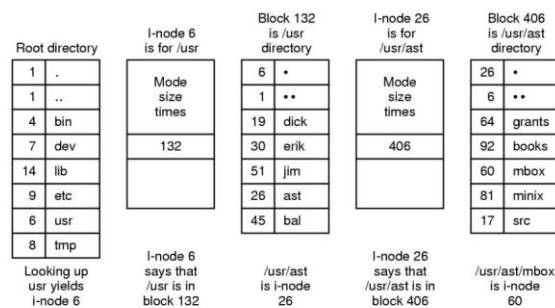
**File System Layout:** File systems usually are stored on disks. Most disks can be divided up into partitions

- Sector 0 of the disk is the **Master Boot Record (MBR)**: The end of the MBR contains the partition table

**Allocation Methods:** How to allocate space on disks for files?

- Files are a sequence of bytes, and disks are arrays of sectors (**512 bytes**)
- **Block Size:** Usually consists of a number of sectors
- File systems view the disk as an array of blocks
- **Contiguous Allocation:** Start and length are stored in the FCB
  - o Easy to implement and excellent read performance
  - o Causes fragmentation and will need periodic compaction (Defrag)
- **Linked List Allocation:** Each file is a linked list of disk blocks. The directory contains a pointer to the first and last block of the file. Each block contains a pointer to the next block and the last block contains a NIL pointer (-1)
  - o No fragmentation
  - o Random access is slow
- **Indexed Allocation:** Brings all pointers together into an index block. Index block contains all indexes of the block

**Entry Lookup:** The **superblock** has the location of the i-node which represents the root directory. Once the root directory is located a search through the directory tree finds the desired directory entry (block)



The steps in looking up `/usr/ast/mbox`

**Free Space Management:** Limited disk space, so we need to reuse the space from deleted files

- **Bitmap:** One bit for each block on the disk. If the block is free the bit is 1 else the bit is 0. Makes it easy to find a contiguous group of free blocks
  - o Used in Apple machines
- Uses Linked List, Indexing, or a combination of the 2

## **NFS:** Network File System

- **NFS File Server:** Runs on a machine (that may have large disks) that has a local file system
- **NFS Client:** Runs on an arbitrary machine and access the files on machines that run NFS servers
- **Mounted:** Means that a directory can appear in the tree structure
- **Export List:** Maintained by server, list of local directories that server exports for mounting and names of machines that are permitted to mount them

**NFS Mount Protocol:** Following a mount request that conforms to its export list, the servers returns a **file handle** (a key for further access)

- The mount operation changes only the user's view and does affect the server side