

Computer Science 3305 – Operating Systems – Course Notes

(Lecture #1 – Course Introduction)

- No material to study from this lecture.

(Lecture #2 – OS Introduction History & Processes – p1 & Forks – p1)

- Operating System = the software layer between user applications and hardware, that manages hardware resources.
- *History of OS:*
 - 1st Gen – vacuum tubes and plug boards (direct input).
 - Ran one calculation at a time.
 - No languages or OS.
 - 2nd Gen – transistors (batch systems).
 - Programs written on paper in FORTRAN or assembly and encoded on punch cards.
 - 3^d Gen – integrated circuits and multiprogramming.
 - Allowed for multiple calculations at the same time.
 - 4th and Current Gen – large scale integration and personal computers.
 - First OS offered by bill gates called DOS.
 - Doug Englehart at Stanford invented the GUI.
 - Jobs designed the first user friendly PC.
 - Next Gen – systems connected by high-speed networks.
- *History of Unix:*
 - Multics = first large timesharing system developed by MIT, General Electric and Bell Labs.
 - Ken Thompson found a computer that nobody was using, and wrote a stripped-down one user version of Multics in C, which became Unix.
 - MINIX = mini-Unix system developed by Andrew Tanenbaum.
 - A Finish student, Linus Torvalds wanted to add stuff until he eventually made his own OS called Linux.
- Process = an instance of a running application. As it runs it changes states.
 - New – process created.
 - Running – instructions being executed.
 - Waiting – process is waiting for an event to occur.
 - Terminated – process has finished execution.

- Process Control Block (PCB) = each process in the OS is represented by a block.
 - Context information would include:
 - Process identifier (PID).
 - Process state.
 - Program counter.
 - CPU registers.
 - CPU-Scheduling information.
 - Memory-Management information.
 - I/O status information.
- Fork() = a system call that creates a child process that is a duplicate of the parent.
 - Created by Unix.
 - Child inherits the state of its parent process, with the same instructions, variables, variable values and position.
 - The parent and child have separate copies of that state though.
 - If the call is SUCCESSFUL it returns the child PID to the parent, and returns 0 to the child.
 - If the call FAILS it returns -1 to the parent, and no child is created.
- Pid_t = data type that represents process identifiers.
 - You can call getpid() – returns the PID of calling process.
 - You can call getppid() – returns the PID of parent process.

(Lecture #3 – Processes – p2 & Forks – p2)

- Processes get a share of the CPU to give other processes a turn. The switching between processes (like parent or child) depends on many factors, such as machine load or process scheduling.
- Nondeterministic = output interleaving, where you can't determine output by looking at code.

- Fork Example #3.c


```
#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <sys/wait.h>
#include <unistd.h>
#include <errno.h>
/*
 * This program forks a process. The child and parent processes print to terminal to identify themselves
 */
int main ()
{
    pid_t i, j, pid;
    pid=fork();
    if (pid <0) // fork unsuccessful
    {
        printf("fork unsuccessful");
        exit(1);
    }
    if (pid>0) //parent
    {
        i = getpid();
        wait(NULL);
        printf("\n I am parent with PID %d\n", i);
    }
    if (pid==0) //child
    {
        i = getpid();
        j = getppid();
        printf("\n I am a child with PID %d and my parent's PID is %d\n", i, j);
    }
}
```

- Fork Example #3.c – possible output:

Possibility #1	Possibility #2
PARENT 0	PARENT 0
PARENT 1	PARENT 1
PARENT 2	PARENT 2
PARENT 3	PARENT 3
PARENT 4	PARENT 4
PARENT 5	PARENT 5
PARENT 6	PARENT 6
PARENT 7	CHILD 0
PARENT 8	CHILD 1
PARENT 9	CHILD 2
CHILD 0	
CHILD 1	
CHILD 2	
CHILD 3	
CHILD 4	
CHILD 5	
CHILD 6	
CHILD 7	
CHILD 8	
CHILD 9	
	PARENT 7
	PARENT 8
	PARENT 9
	CHILD 3
	CHILD 4
	CHILD 5
	CHILD 6
	CHILD 7
	CHILD 8
	CHILD 9

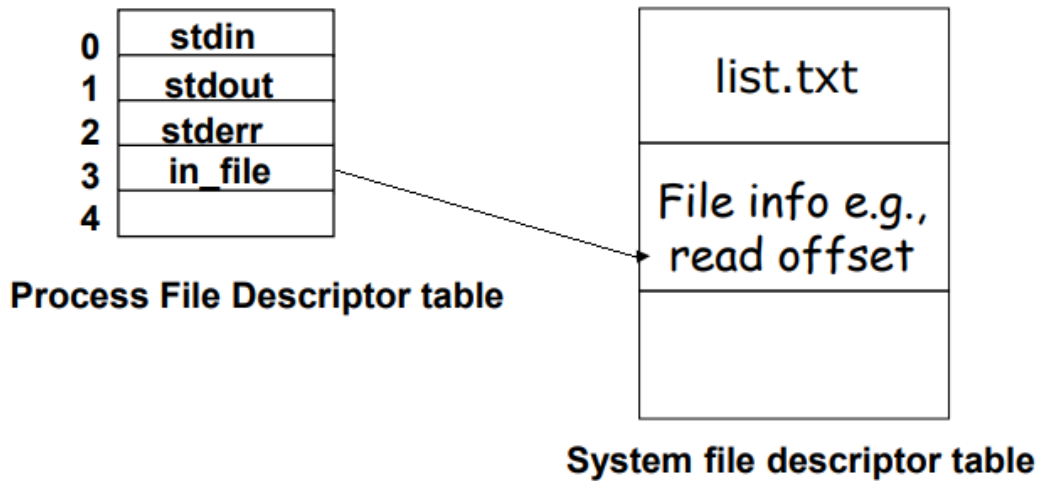
- And many more possibilities as well!
- Amount of children created by forks = $(2^{\text{\#forks}}) - 1$.

(Lecture #4 – Processes – p3 & Forks – p3)

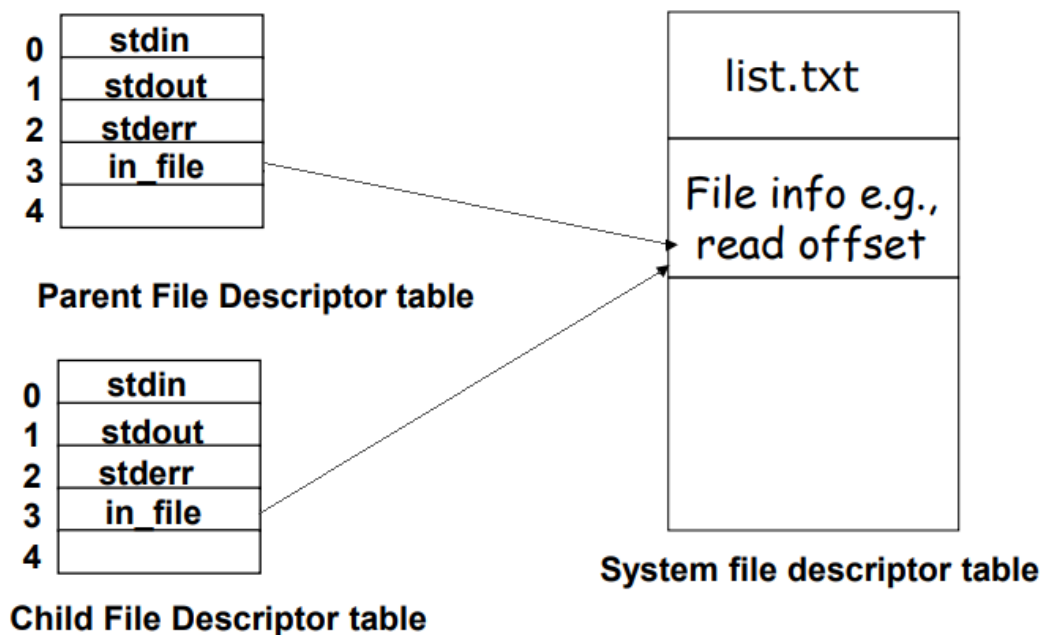
- Process File Descriptor Table = held by every process, where each entry represents something that can be read or written from (like stdin or stdout).
- System File Descriptor Table = every open file has one.

Example:

```
int in_file;  
in_file = open("list.txt", O_RDONLY);
```



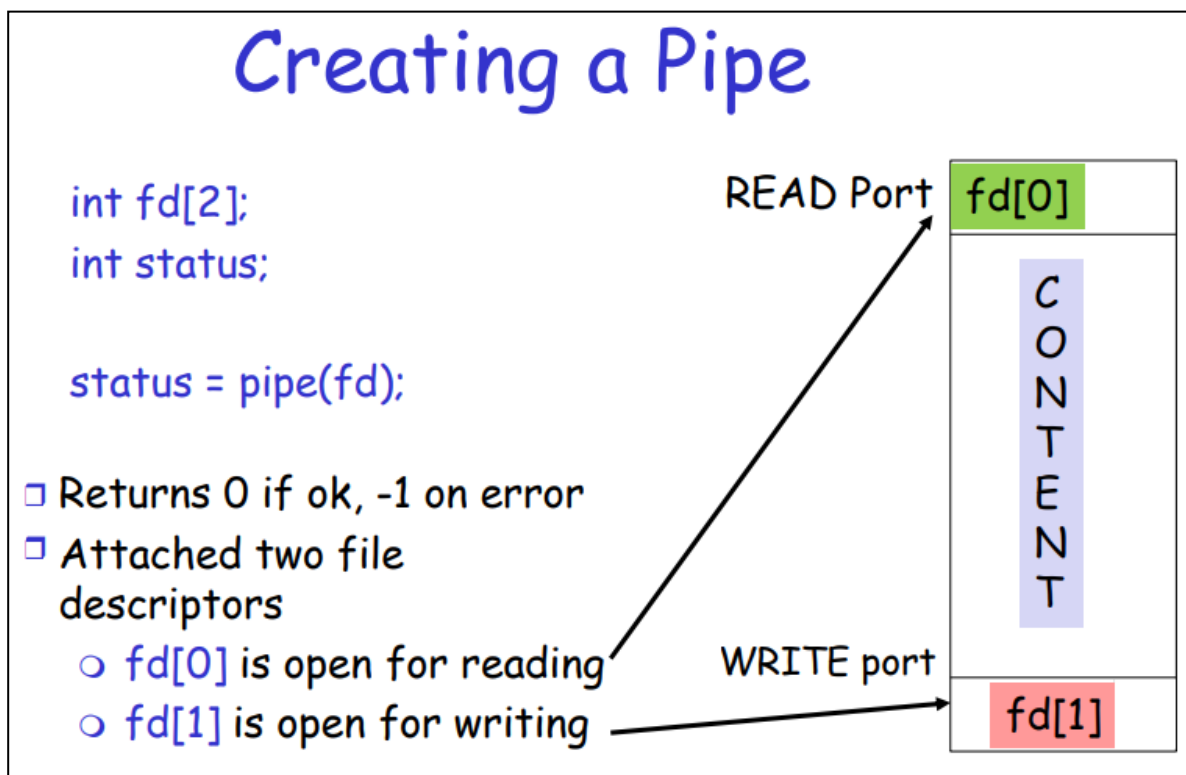
- Fork() = copies the process file descriptor table, meaning the parent and child point to the same entry in the system file descriptor table.



- If you open a file, then use `fork()` → the child process gets a copy of its parent's process file descriptor table, and the parent and child share a file pointer.
 - So if you read in "cat", the first char as parent reads is "c" and the child then reads "a".
- If you use `fork()` then open a file → the child and the parent process each open a file, and have their own entries in the system file descriptor table, and their own pointers.
 - So if you read in "cat", the first char as parent reads "c" and the child reads "c" as well.

- Recommended to open files after forks.
- `execl()` = replaces a process with a new binary file loaded into memory. Only returns (-1) if it fails.
 - The Child will not return to the old program unless `execl` fails.
 - Specifically it stops running the parent (as it's destroyed), and removes the PCB. So the code before `execl()` is run, then `execl()`, and nothing else runs after it.

(Lecture #5 – Processes – p4 & Pipes)



- Pipe won't allow you to read an empty pipe or write to a full pipe.
- `Read(fd[0], &content, byte size);`
 - `printf("\n from child: this is what I read %c", content);`
- `Write(fd[1], content to write, byte size);`

(Lecture #6 – Signals & System Calls)

- Signal = notifies a process that an event has occurred, and the normal execution is interrupted.
- Signal() = captures a specific event and associates it with a programmer-defined function.
 - However you must include signal.h for this.
- SIGINT = interrupt signal from terminal (control-c).
- SIGALRM = instructs the OS to send an alarm.
 - Signal (SIGALRM, alarmHandler_Variable);
 - Alarm(3);
- All OS have an interface that is accessible by users/user programs and GUI's.
- System calls = provide an interface to OS services by passing relevant information to it.
- Application Programmer Interface (API) = where the programmer only needs to understand the system function through its parameters and results.
- System Function = used by the user / API programmer.
 - Some for process management (like pid or exit).
 - Some for file management (like fd, read, write).
 - Some for directory management (like mkdir or rmdir).
- System Call = part of the OS Kernel.
- System calls pass parameters as registers, blocks (used by Linux and Solaris) or stacks.
- Linux system calls with fewer than 6 parameters are passed in registers. Otherwise as blocks.
- The OS maintains a system call handler table which is indexed according to the system call numbers.
- TRAP = switches CPU to supervisor (kernel mode).
 - This saves the state of the user process so that the OS instructions can be executed and then the user process can execute after.
- What happens when you make a system call (in Linux)?
 - Step 1 – the input parameters are passed into register or to a block.
 - Step 2 – TRAP is executed.
 - Saving the user process as T.
 - System call number for the specific function is sent to the system call handler.

- This call number tells the OS what system call handler (kernel code) to execute.
- This causes a switch from the user mode to the kernel mode.
- Step 3 – system call handler code executes (sometimes will have to wait for data from the disk to be slowly read in).
- Step 4 – after execution, control returns to the system function.

(Lecture #7 – Multiprogramming & Threading)

- Multiprogramming = allows for the execution of multiple processes, but only one active at a time.
- Interleaved Execution = allows for a process to be running at all times in order to maximize CPU utilization.
- Context Switching = when the OS suspends the current process so that another process can execute.
 - The OS must capture information about the process before it switches, so that it can restore the hardware to the same configuration after.
- Process Control Block (PCB) = represents each process in the OS.
 - Contains all of the states for a program, including (but not limited to):
 - Pointers.
 - Program counter (PC).
 - Current values.
 - Operating resources (open files and network connections).
 - Process identifier (PID).
 - Process priority.
- Process Execution States = an instance of a running application. As it runs it changes states.
 - New – process is being created.
 - Ready – process is waiting to be assigned to a processor.
 - Running – instructions being executed.
 - Waiting – process is waiting for an event to occur.
 - Exit – process has finished execution.
 - Note: only one process can be running on a processor at an instant, but many can be ready or waiting.
- Job Queue = contains all processes in the system, with each new process entering the system coming here.
- Ready Queue = contains all processes ready and waiting to be executed.
- Queues are implemented using a linked list.

- Fork is not efficient for multiprocessing as it copies everything.
- Threads:
 - Share code, data, open files and sockets.
 - But, they have their own CPU context with a program counter (PC), stack pointer (SP) and register state.
- Pthread Library = common thread library, implemented with pthread.h.
 - Each thread has a unique ID – called with pthread_self().
 - Thread ID's are of type pthread_t (typically an unsigned int).

pthread_create()

Creates a new thread

```
int pthread_create (
    pthread_t *thread,
    pthread_attr_t *attr,
    void * (*start_routine),
    void *arg);
```

- Returns 0 to indicate success, otherwise returns error code
- **thread**: name of the new thread
- **attr**: argument that specifies the attributes of the thread to be created (NULL = default attributes)
- **start_routine**: function to use as the start of the new thread
- **arg**: argument to pass to the new thread routine

pthread_create() example

Let us say that you want to create a thread that simply prints "hello world...I am a thread"

```
int main(int argc, char *argv) {
    pthread_t worker_thread;

    if (pthread_create(&worker_thread, NULL, do_work, NULL) {
        printf("Error while creating thread\n");
        exit(1);
    }
    ...
}

void *do_work() {
    Printf ("\n hello world..I am a thread");
    return NULL;
}
```


- Note: sharing global variables is dangerous as 2 threads may attempt to modify the same variable at the same time.
 - To solve this, use support for mutual exclusion primitives that can be used to protect against this problem.
 - The idea is to lock something before accessing any global variables, and to unlock when done.


(Lecture #8 – CPU Scheduling – p1)

- Process Execution = consists of CPU execution and I/O wait/burst time.
- CPU/I/O Burst Cycle = alternating between CPU bursts and I/O bursts, where the CPU must wait for I/O processes to occur in between running (which is the limited factor as the CPU is fast).
- CPU Scheduler = selects a process from the ready queue.
 - Non-preemptive process:
 - Runs until it relinquishes CPU control.
 - Preemptive process:
 - Runs for a maximum fixed time (set by a clock interrupt at the end time interval).
- Scheduling Evaluation Metrics = quantitative criteria for evaluating a scheduling algorithm.
 - CPU Utilization = percentage of time the CPU is being used.
 - Throughput = the number of completed process per time unit. #Cp's/T.
 - Waiting Time = time spent on the ready queue.
 - Turnaround Time = time from submission till completion of process.
 - Fairness = no process suffers starvation.
- *Scheduler Options:*
 - First Come, First Served (FCFS).
 - Last In First Out (LIFO).
 - Shortest Job First.
 - Priority Scheduling.
 - Round Robin (RR).
 - May use priorities to determine who runs next.
 - Multilevel Queuing.
 - Multilevel Queuing with Feedback.
 - Lottery Scheduling.
- First Come, First Served (FCFS) = the process that requests the CPU first is then allocated the CPU first.
 - Simple to write and understand.
 - However if the first process is the longest, the average wait time is far higher.

(FCFS) Scheduling Example

Process	CPU Burst Time
P1	24
P2	3
P3	3

Assume processes serving request in the order: P1, P2, P3. The scheduling chart:




Waiting time for $P_1 = 0$; $P_2 = 24$; $P_3 = 27$

Average waiting time: $(0 + 24 + 27)/3 = 17$

Instead:

Suppose that the processes serving request in the order P_2, P_3, P_1

The Gantt chart for the schedule is:



Waiting time for $P_1 = 6$; $P_2 = 0$; $P_3 = 3$

Average waiting time: $(6 + 0 + 3)/3 = 3$

Much better than previous case

- Last-In First-Out (LIFO):
 - New processes are placed at the head of the queue.
 - However can lead to starvation.
 - Starvation = early process unable to ever get CPU use.
- Shortest Job First (SJF):
 - Estimated CPU burst time associated with each process.
 - Optimal process that gives minimum average waiting time for a given set of processes.

- Priority Scheduling:
 - A priority number (integer) is associated with each process.
 - The CPU then allocates the process with the highest priority (preemptive, or non-preemptive).
 - Problem – Starvation.
 - Solution – Aging.
 - Aging = as time progresses, increase the priority of the process.
- Round Robin (RR):
 - Each process gets a small unit of CPU time (time quantum (such as 10-100 milliseconds)), and after the time has elapsed the process is paused and added to the end of the ready queue.
 - With 'n' processes and 'q' time quantum, each process gets at most q time units at once.
 - Maximum wait time for the nth process in the ready queue = (n-1)*q.
 - If 'q' is too large → RR becomes like FIFO.
 - If 'q' is too small → RR gets high context switching.
- Multilevel Queuing:
 - Used by most schedulers currently.
 - Makes the ready queue into multiple separate queues by classifying them based on scheduling.
 - Every queue has its own scheduling algorithm, such as RR with q of 5 or 8, or FIFO...etc.
 - Fixed priority scheduling occurs where all processes from a specific queue are executed before moving on to the next queue.
 - Problem → has a possibility of starvation.
- Multilevel Queuing with Feedback:
 - A process is able to move between queues.
 - Processes are separated according to their CPU bursts.
 - If a process uses too much CPU time it is moved to a lower-priority queue.
 - Includes aging for process in lower-priority queue.
- Lottery Scheduling:
 - Each process is given a random amount of lottery tickets.
 - To select the next process to run, the scheduler randomly selects a ticket and that process runs.
 - Solves starvation problem.